

Covid -19 diagnosis with 3D CNN and Lung -CT segmentation



Fardeen Hasib Mozumder

ID : 1018062288



Introduction

- ▶ CT can play a vital role in diagnosis of covid-19

Advantages:

- ▶ Rt-PCR has low accuracy (60-70)% and requires some time, whereas diagnosis from CT scan is faster, which is crucial to manage the pandemic.

Cons:

- ▶ Features are not well distinguished from CAP (community acquired pneumonia)

Features:

- ▶ Bilateral ground glass opacities and consolidation



Data

- Using the default dataset SPGC covid dataset
- The dataset is highly non-uniform
- Very high number of Covid positive samples (170 samples) in opposite to non covid samples (CAP : 60 samples or Healthy : 76 samples)
- Structural difference of images for different patients (ie : difference in number of slices, Hounsfield Unit range, Pixel spacing, slice thickness)



Data Preprocessing

Data Splitting :

- ▶ 55 number of samples from each category for training process(40 for training & 15 for validation)
- ▶ 5 from each category for testing
- ▶ (So, total 60 from each category, as CAP has minimum number of samples of 60)

Tackling non-uniformity in Data:

- ▶ Hu range was set to -1000 to 400 (as, air hu \approx -1000 and lung tissue hu \approx 300 addition 100 accounts for infection and blood clots), considering , values outside the range are not of interest diagnosing covid.



Data Preprocessing

Data Normalization:

- Hounsfield Unit normalized to 0 to 1.

[values closer to 0 represent air or non infected region, values closer to 1 represent infected region]

- Lungs were segmented and cropped to reduce loss (details in later slides)
- 3D images was resized to (50,128,128) before passing to training



Methodology:

3D CNN

- Preprocessed data is passed to a 3D CNN model with (4 segments of of 2 3D CNN models followed by a 3D maxpooling layer)
- The kernel size is of 3d CNN layer is (3,3,3) and the activation function is relu
- The pooling size is (2,2,2,) for the maxpooling layer.
- After these, there are 2 dense layers with node of 32 and 3 .
- Activation function of last dense layer is softmax as there are 3 classes.
- A dropout layer is used between the dense layer to avoid overfitting

```
def get_model(width=128, height=128, depth=64):
    """Build a 3D convolutional neural network model."""

    inputs = keras.Input((width, height, depth, 1))

    x = layers.Conv3D(filters=64, kernel_size=3, activation="relu")(inputs)
    x = layers.MaxPool3D(pool_size=2)(x)
    x = layers.BatchNormalization()(x)

    x = layers.Conv3D(filters=64, kernel_size=3, activation="relu")(x)
    x = layers.MaxPool3D(pool_size=2)(x)
    x = layers.BatchNormalization()(x)

    x = layers.Conv3D(filters=128, kernel_size=3, activation="relu")(x)
    x = layers.MaxPool3D(pool_size=2)(x)
    x = layers.BatchNormalization()(x)

    x = layers.Conv3D(filters=256, kernel_size=3, activation="relu")(x)
    x = layers.MaxPool3D(pool_size=2)(x)
    x = layers.BatchNormalization()(x)

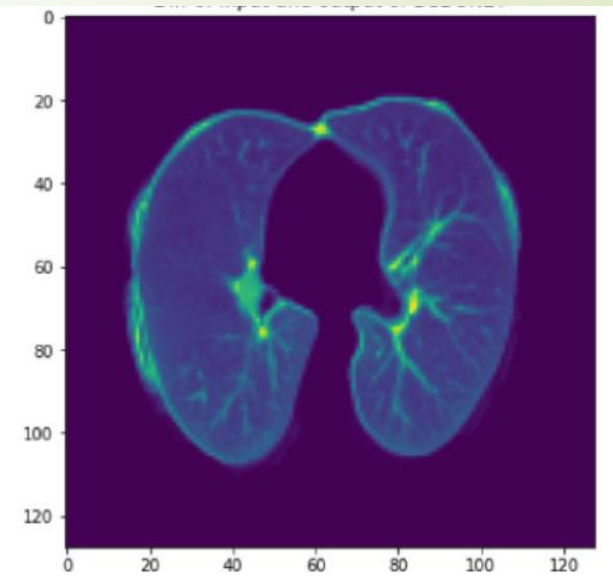
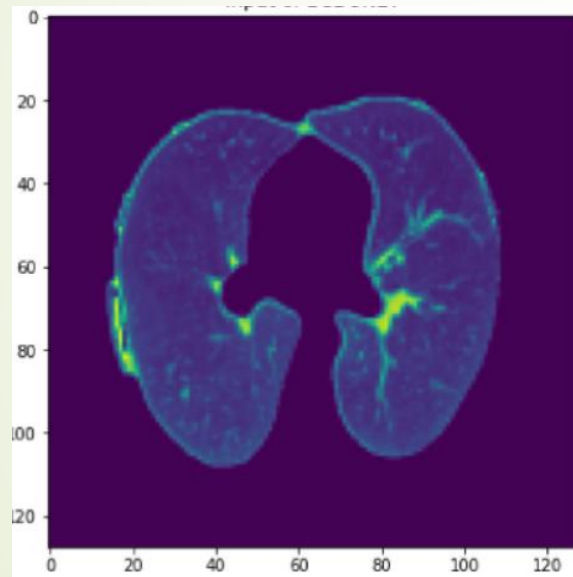
    x = layers.GlobalAveragePooling3D()(x)
    x = layers.Dense(units=512, activation="relu")(x)
    x = layers.Dropout(0.3)(x)

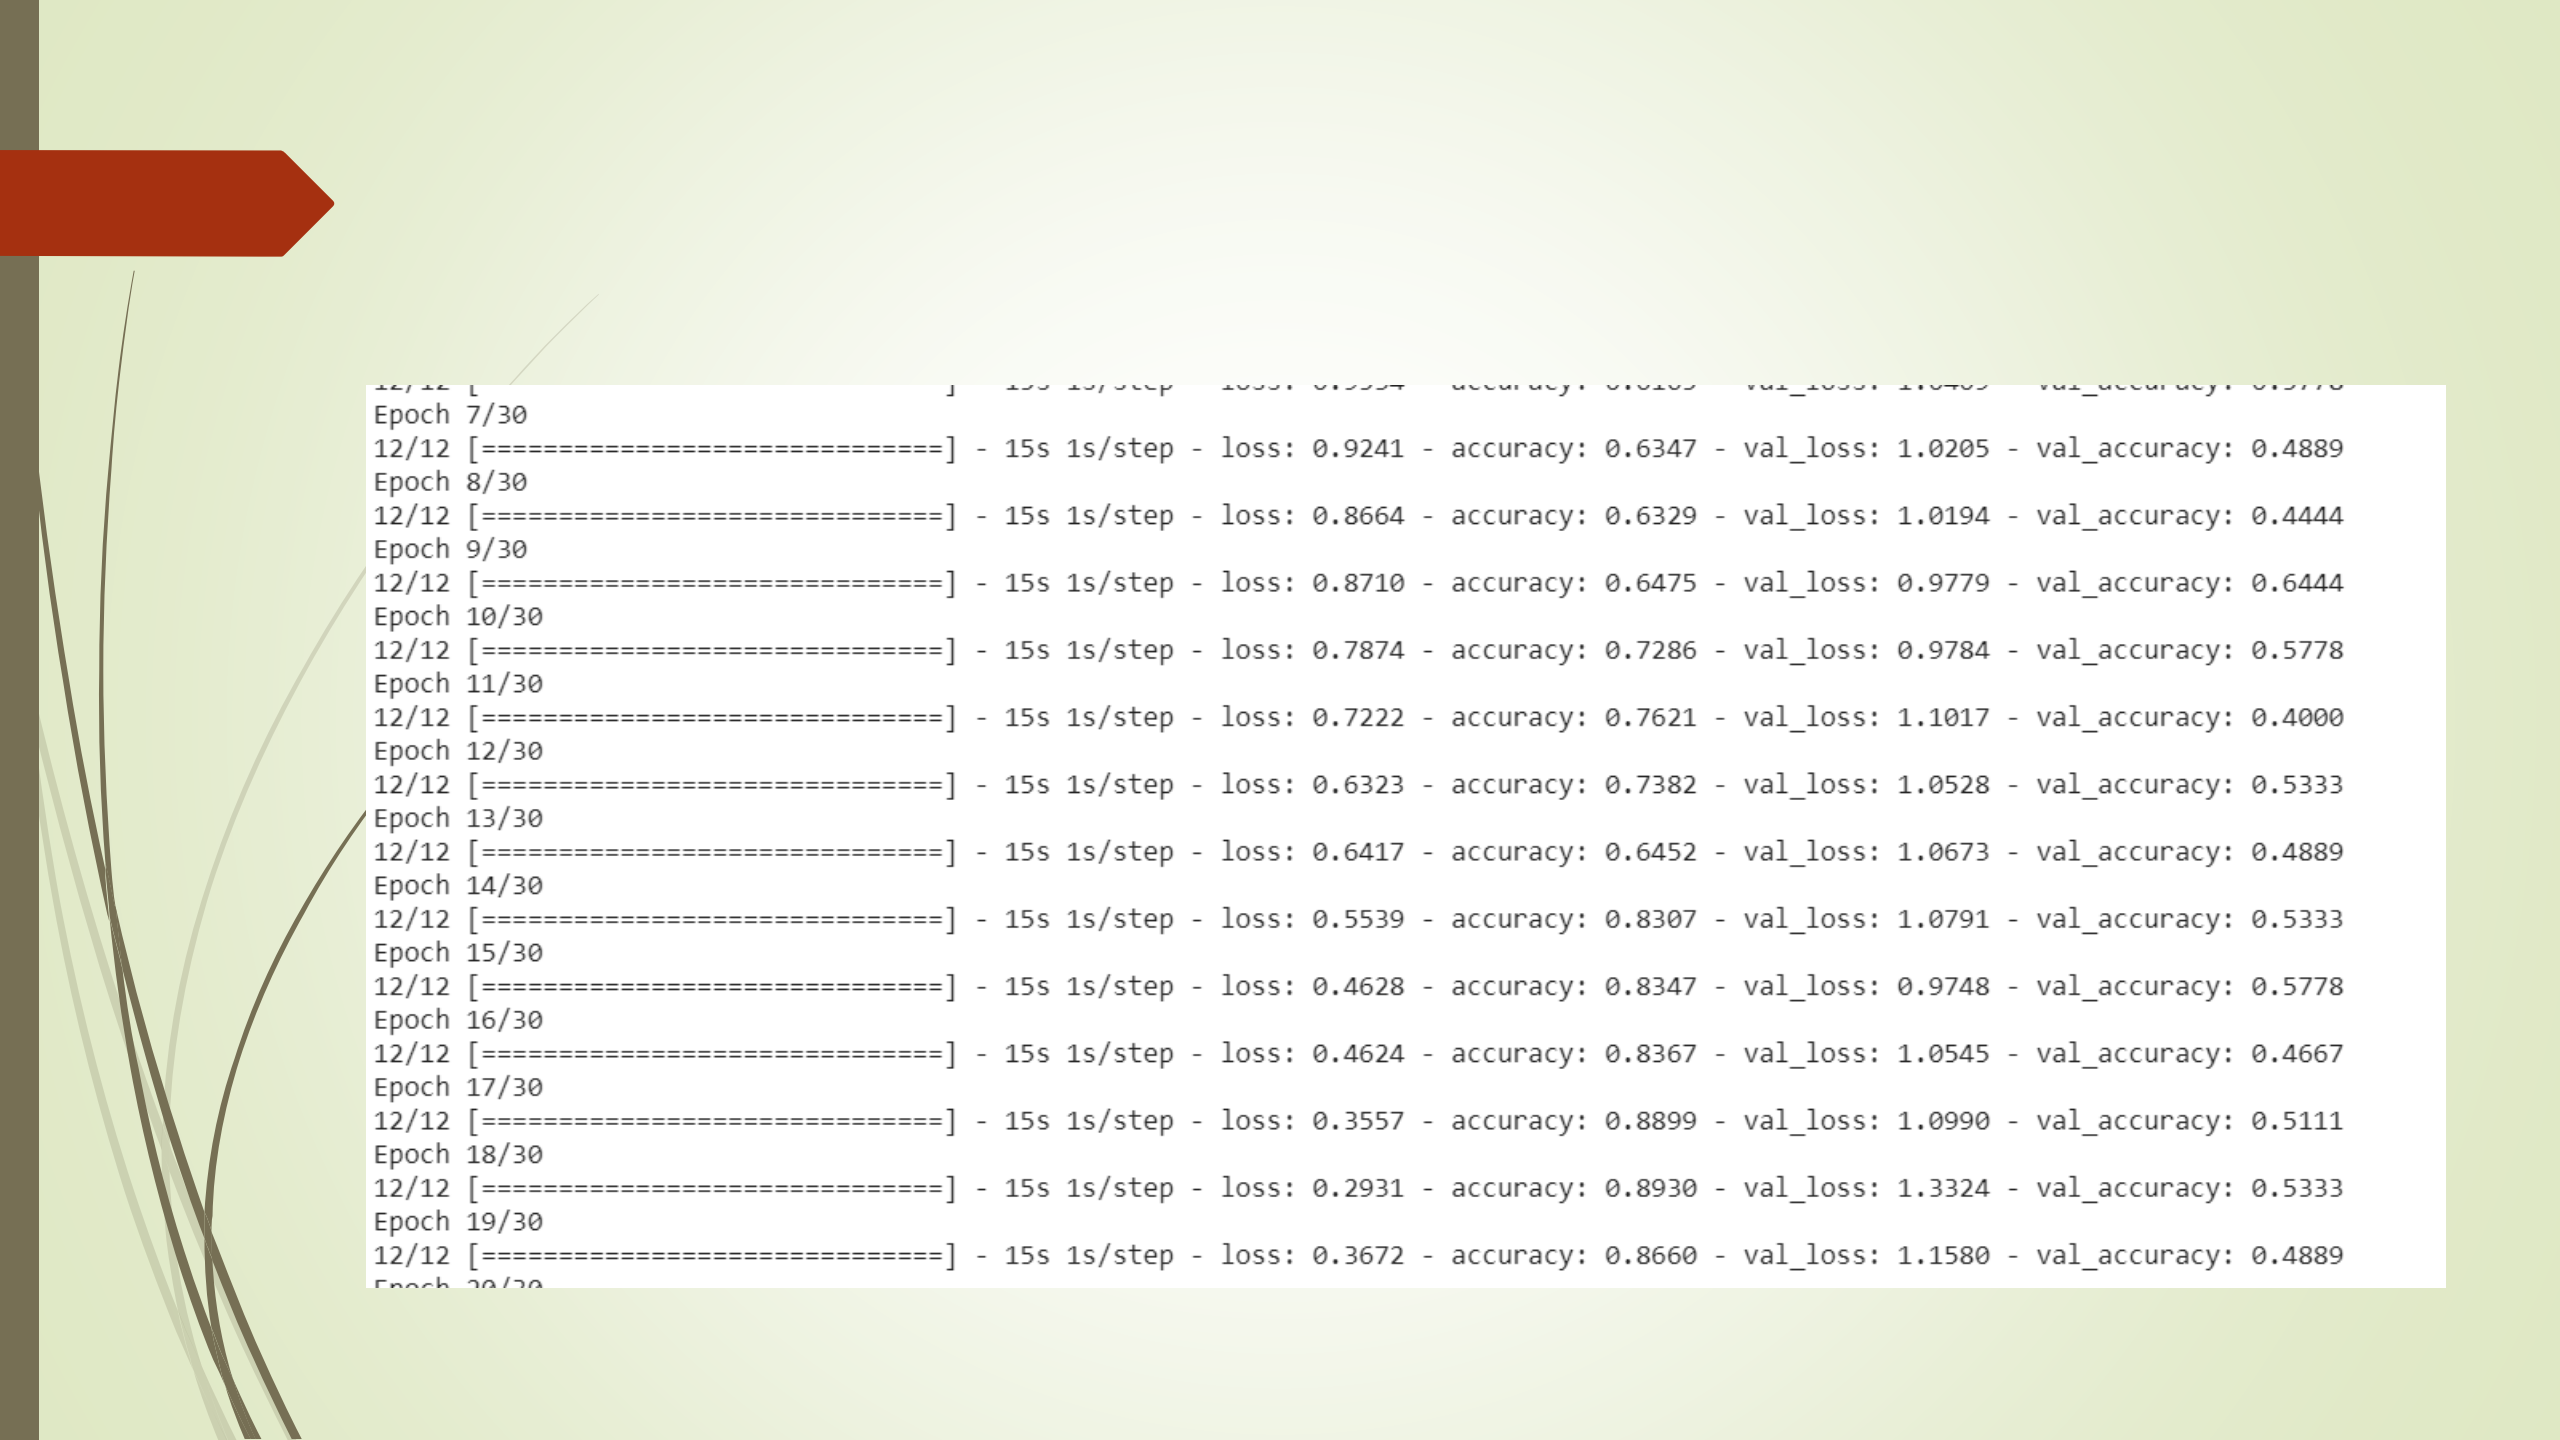
    outputs = layers.Dense(units=1, activation="sigmoid")(x)

    # Define the model.
    model = keras.Model(inputs, outputs, name="3dcnn")
    return model

# Build model.
model = get_model(width=128, height=128, depth=64)
model.summary()
```

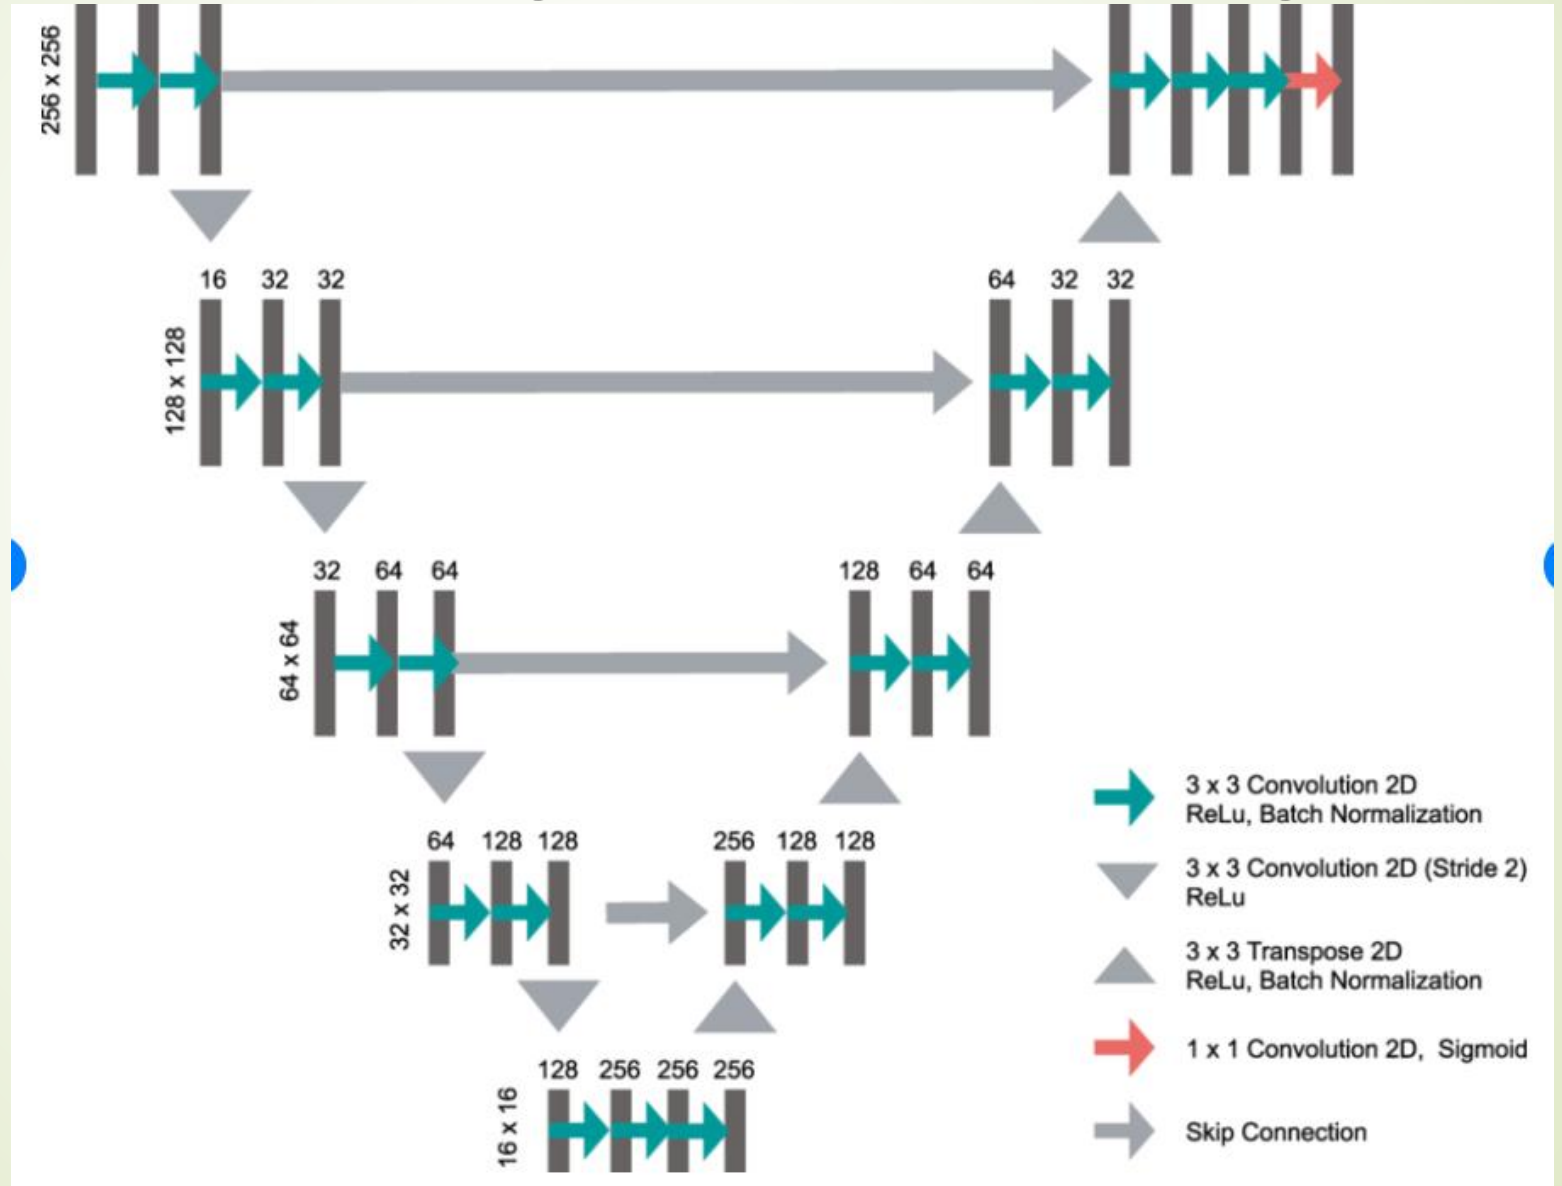

Methodology: 3D CNN



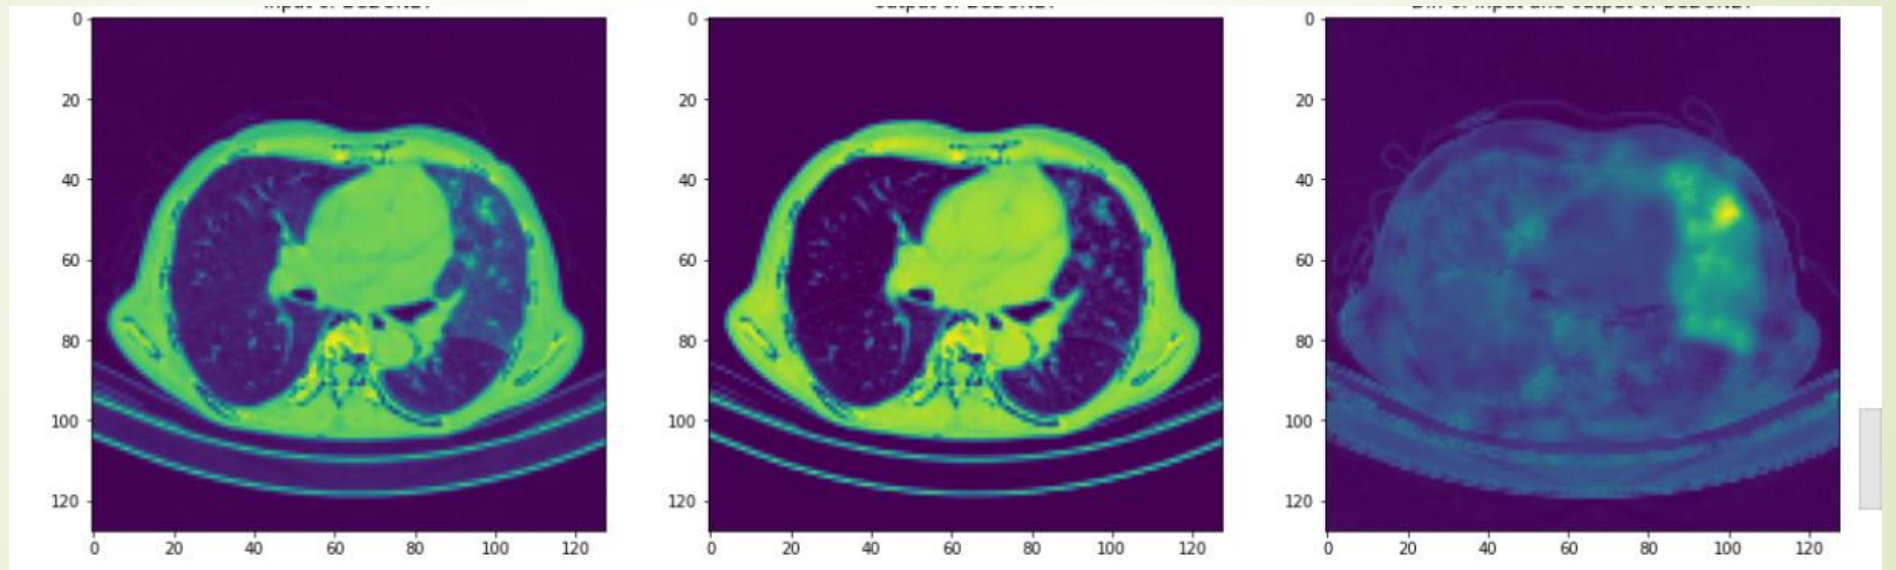


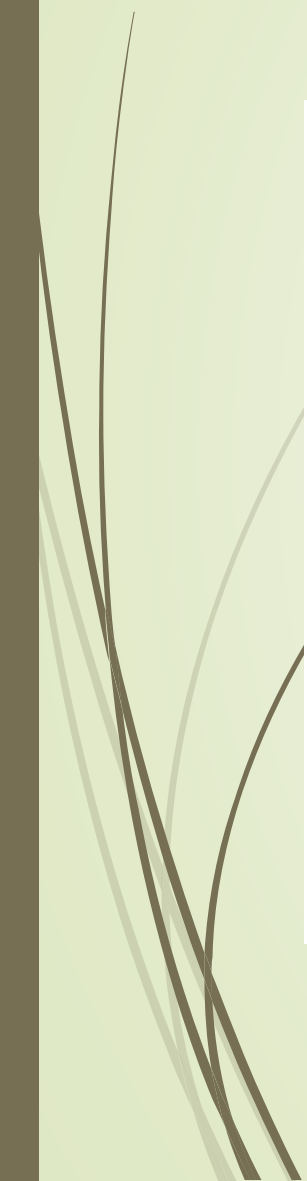

```
12/12 [=====] - 15s 1s/step - loss: 0.9241 - accuracy: 0.6347 - val_loss: 1.0205 - val_accuracy: 0.4889
Epoch 7/30
12/12 [=====] - 15s 1s/step - loss: 0.8664 - accuracy: 0.6329 - val_loss: 1.0194 - val_accuracy: 0.4444
Epoch 8/30
12/12 [=====] - 15s 1s/step - loss: 0.8710 - accuracy: 0.6475 - val_loss: 0.9779 - val_accuracy: 0.6444
Epoch 9/30
12/12 [=====] - 15s 1s/step - loss: 0.7874 - accuracy: 0.7286 - val_loss: 0.9784 - val_accuracy: 0.5778
Epoch 10/30
12/12 [=====] - 15s 1s/step - loss: 0.7222 - accuracy: 0.7621 - val_loss: 1.1017 - val_accuracy: 0.4000
Epoch 11/30
12/12 [=====] - 15s 1s/step - loss: 0.6323 - accuracy: 0.7382 - val_loss: 1.0528 - val_accuracy: 0.5333
Epoch 12/30
12/12 [=====] - 15s 1s/step - loss: 0.6417 - accuracy: 0.6452 - val_loss: 1.0673 - val_accuracy: 0.4889
Epoch 13/30
12/12 [=====] - 15s 1s/step - loss: 0.5539 - accuracy: 0.8307 - val_loss: 1.0791 - val_accuracy: 0.5333
Epoch 14/30
12/12 [=====] - 15s 1s/step - loss: 0.4628 - accuracy: 0.8347 - val_loss: 0.9748 - val_accuracy: 0.5778
Epoch 15/30
12/12 [=====] - 15s 1s/step - loss: 0.4624 - accuracy: 0.8367 - val_loss: 1.0545 - val_accuracy: 0.4667
Epoch 16/30
12/12 [=====] - 15s 1s/step - loss: 0.3557 - accuracy: 0.8899 - val_loss: 1.0990 - val_accuracy: 0.5111
Epoch 17/30
12/12 [=====] - 15s 1s/step - loss: 0.2931 - accuracy: 0.8930 - val_loss: 1.3324 - val_accuracy: 0.5333
Epoch 18/30
12/12 [=====] - 15s 1s/step - loss: 0.3672 - accuracy: 0.8660 - val_loss: 1.1580 - val_accuracy: 0.4889
Epoch 19/30
```

Automatic segmentation using BCDU



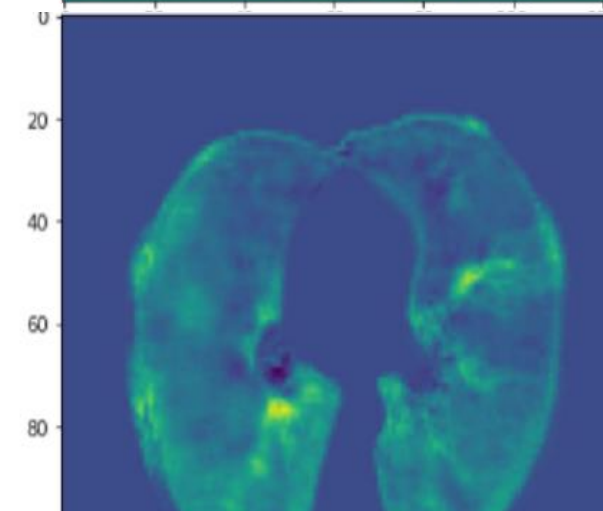
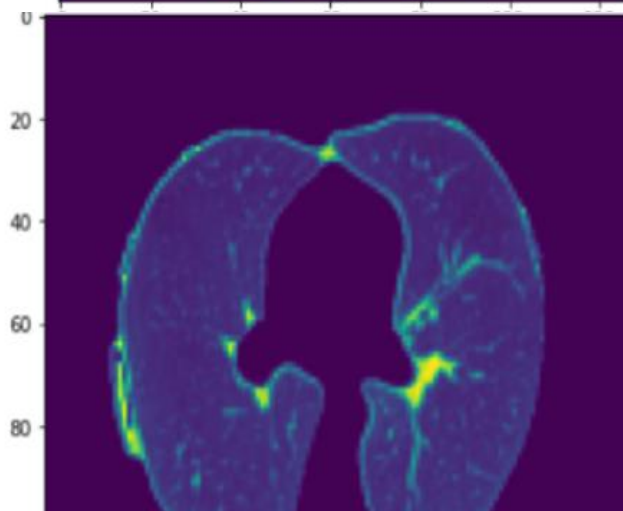
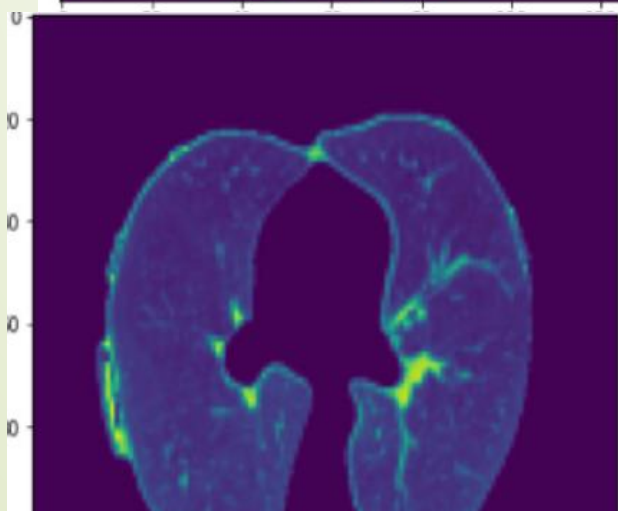
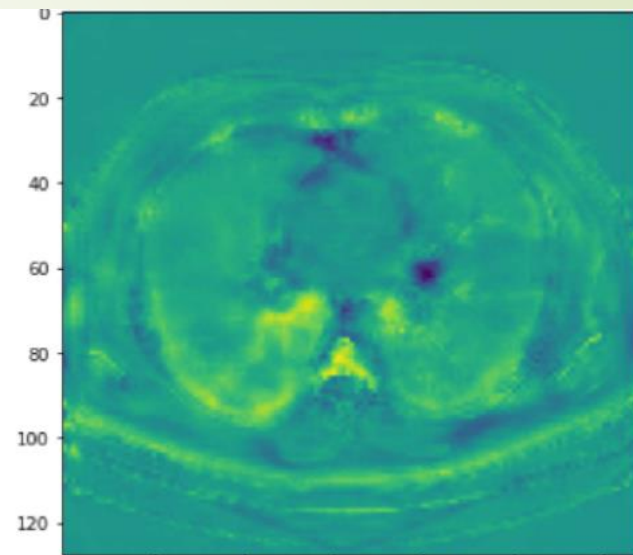
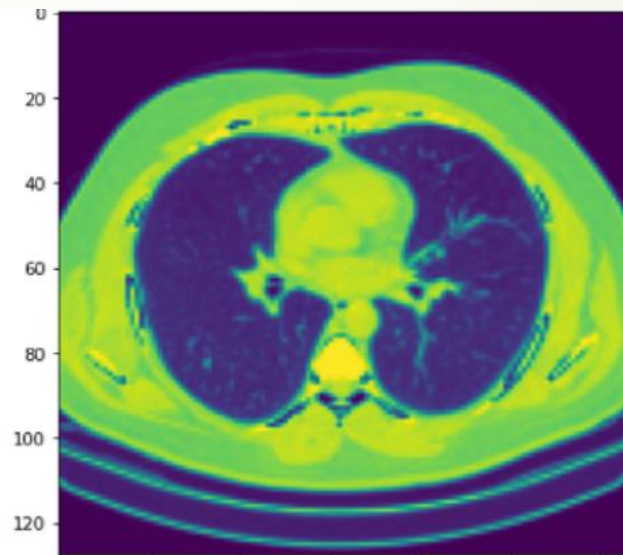
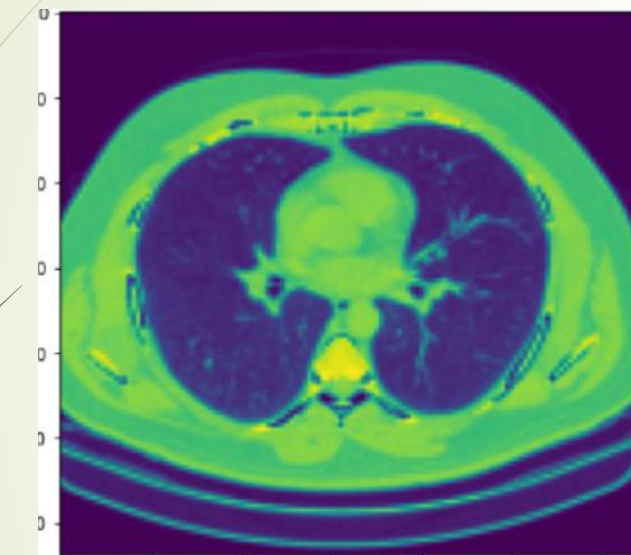
Effects of BCDU





```
Epoch 7/30
12/12 [=====] - 5s 450ms/step - loss: 0.8202 - accuracy: 0.5815 - val_loss: 1.0235 - val_accuracy: 0.5333
Epoch 8/30
12/12 [=====] - 5s 449ms/step - loss: 0.7624 - accuracy: 0.7124 - val_loss: 1.1268 - val_accuracy: 0.5778
Epoch 9/30
12/12 [=====] - 5s 450ms/step - loss: 0.6727 - accuracy: 0.6712 - val_loss: 1.1578 - val_accuracy: 0.6000
Epoch 10/30
12/12 [=====] - 5s 454ms/step - loss: 0.7072 - accuracy: 0.7449 - val_loss: 1.0080 - val_accuracy: 0.6444
Epoch 11/30
12/12 [=====] - 5s 452ms/step - loss: 0.6802 - accuracy: 0.6724 - val_loss: 1.0460 - val_accuracy: 0.5778
Epoch 12/30
12/12 [=====] - 5s 455ms/step - loss: 0.6100 - accuracy: 0.7715 - val_loss: 1.0732 - val_accuracy: 0.5333
Epoch 13/30
12/12 [=====] - 5s 455ms/step - loss: 0.5241 - accuracy: 0.8066 - val_loss: 1.1776 - val_accuracy: 0.6889
Epoch 14/30
12/12 [=====] - 5s 453ms/step - loss: 0.4626 - accuracy: 0.8258 - val_loss: 1.2523 - val_accuracy: 0.6000
Epoch 15/30
12/12 [=====] - 5s 457ms/step - loss: 0.3766 - accuracy: 0.8716 - val_loss: 1.1472 - val_accuracy: 0.6000
Epoch 16/30
12/12 [=====] - 5s 456ms/step - loss: 0.3653 - accuracy: 0.8662 - val_loss: 1.5842 - val_accuracy: 0.6000
Epoch 17/30
12/12 [=====] - 5s 455ms/step - loss: 0.6362 - accuracy: 0.7275 - val_loss: 1.3897 - val_accuracy: 0.6000
Epoch 18/30
12/12 [=====] - 5s 456ms/step - loss: 0.4058 - accuracy: 0.8228 - val_loss: 1.2510 - val_accuracy: 0.6222
Epoch 19/30
```


Effects of preprocessed lung cropping and segmentation



Effects of preprocessed lung cropping and segmentation

```
Epoch 1/3  
3290/3290 [=====] - 897s 260ms/step - loss: 0.3966  
Epoch 2/3  
3290/3290 [=====] - 850s 258ms/step - loss: 0.3727  
Epoch 3/3  
3290/3290 [=====] - 846s 257ms/step - loss: 0.3722
```

```
3290/3290 [=====] - 318s 97ms/step - loss: 0.0760  
Epoch 3/3  
3290/3290 [=====] - 318s 97ms/step - loss: 0.0752
```

Testing on uncropped Lungs with BCDU net

```
[ ] predicted = loaded_model.predict(dataset)

Label = ['Control', 'COVID-19', 'CAP']
for i in range(predicted.shape[0]):
    lbl = np.argmax(predicted[i])
    print('The case number %d is %i, test2[i], 'predicted', lbl, 'with probability of %.2f'%(100*predicted[i, lbl]))
```

```
The case number 0 is 1 predicted 1 with probability of 99.96
The case number 1 is 1 predicted 1 with probability of 97.25
The case number 2 is 1 predicted 0 with probability of 99.96
The case number 3 is 1 predicted 0 with probability of 84.47
The case number 4 is 1 predicted 1 with probability of 71.83
The case number 5 is 0 predicted 0 with probability of 99.77
The case number 6 is 0 predicted 0 with probability of 99.98
The case number 7 is 0 predicted 0 with probability of 100.00
The case number 8 is 0 predicted 0 with probability of 100.00
The case number 9 is 0 predicted 0 with probability of 99.91
The case number 10 is 2 predicted 2 with probability of 87.92
The case number 11 is 2 predicted 2 with probability of 99.68
The case number 12 is 2 predicted 1 with probability of 88.25
The case number 13 is 2 predicted 2 with probability of 59.80
The case number 14 is 2 predicted 1 with probability of 92.08
```


Testing on cropped Lungs with BCDU net

```
Label = ['Control', 'COVID-19', 'CAP']  
for i in range(predicted.shape[0]):  
    lbl = np.argmax(predicted[i])  
    print('The case number %d is '%i, test2[i], 'predicted', lbl, 'with probability
```

```
The case number 0 is 1 predicted 1 with probability of 100.00  
The case number 1 is 1 predicted 1 with probability of 98.07  
The case number 2 is 1 predicted 1 with probability of 94.27  
The case number 3 is 1 predicted 1 with probability of 87.93  
The case number 4 is 1 predicted 2 with probability of 99.93  
The case number 5 is 0 predicted 0 with probability of 60.45  
The case number 6 is 0 predicted 0 with probability of 100.00  
The case number 7 is 0 predicted 0 with probability of 100.00  
The case number 8 is 0 predicted 0 with probability of 98.36  
The case number 9 is 0 predicted 0 with probability of 99.94  
The case number 10 is 2 predicted 2 with probability of 99.92  
The case number 11 is 2 predicted 2 with probability of 100.00  
The case number 12 is 2 predicted 0 with probability of 78.86  
The case number 13 is 2 predicted 1 with probability of 99.99  
The case number 14 is 2 predicted 1 with probability of 100.00
```



Findings

- Segmentation with BCDU net on cropped lungs results in better and more effective diagnosis as it gives fewer false negatives.
- Needs better and more effective segmentation process to differentiate between CAP vs Covid to increase accuracy



Future work

- Exploring more segmentation networks (Unet++, Segnet, Linknet, PSPnet etc)
- Exploring more 3D CNN networks suitable for our data.
- Optimizing the threshold levels of HU to increase efficiency
- Resizing our data to bigger unit before passing to networks



Resource Limitations

- Hardware limitations:
(GPU, RAM)
- Data Limitation :
Small Dataset,
Absence of manually segmented/masked samples to use as ground truth

